A Tour of TensorFlow

Peter Goldsborough

May 31, 2016

A Tour of TensorFlow

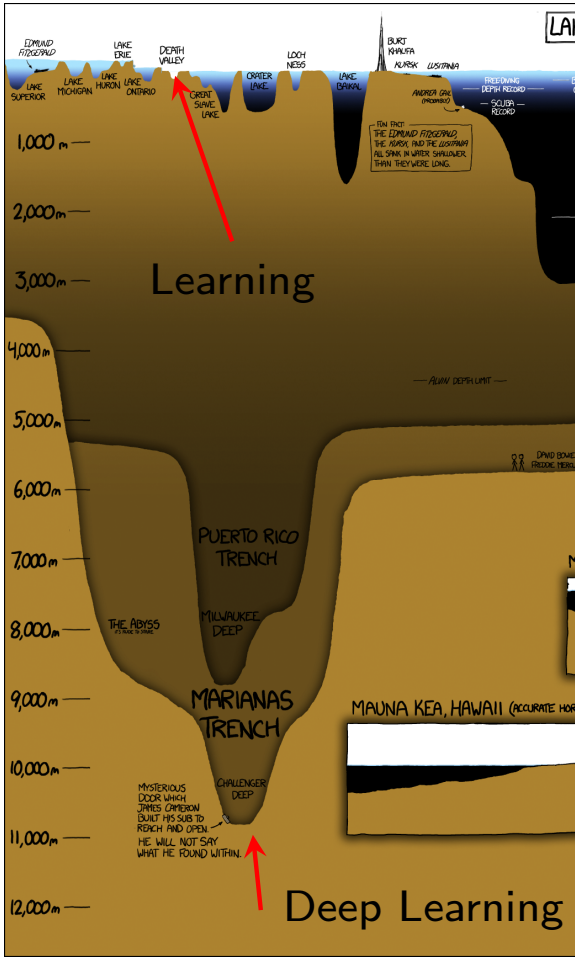

A Tour of TensorFlow

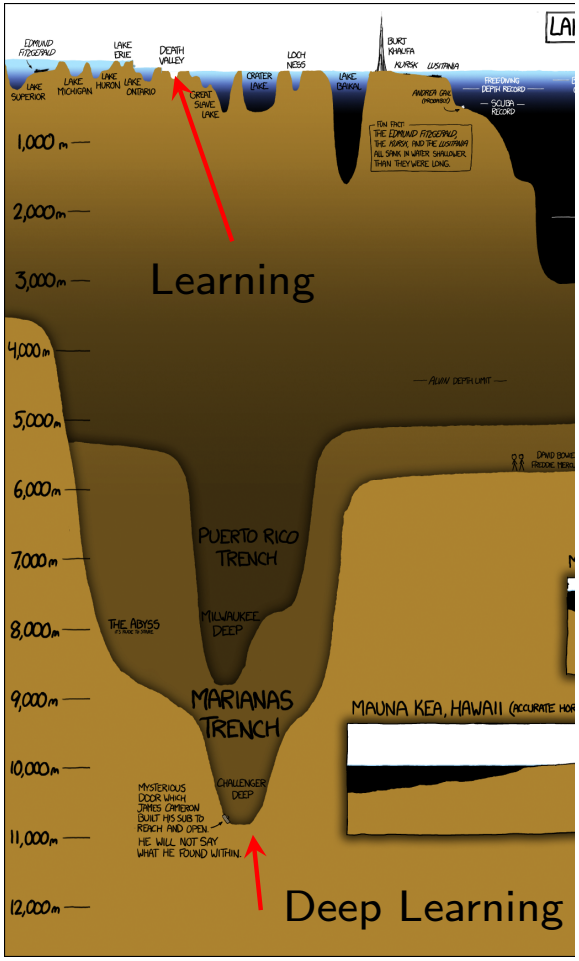

TensorFlow is

A Tour of TensorFlow

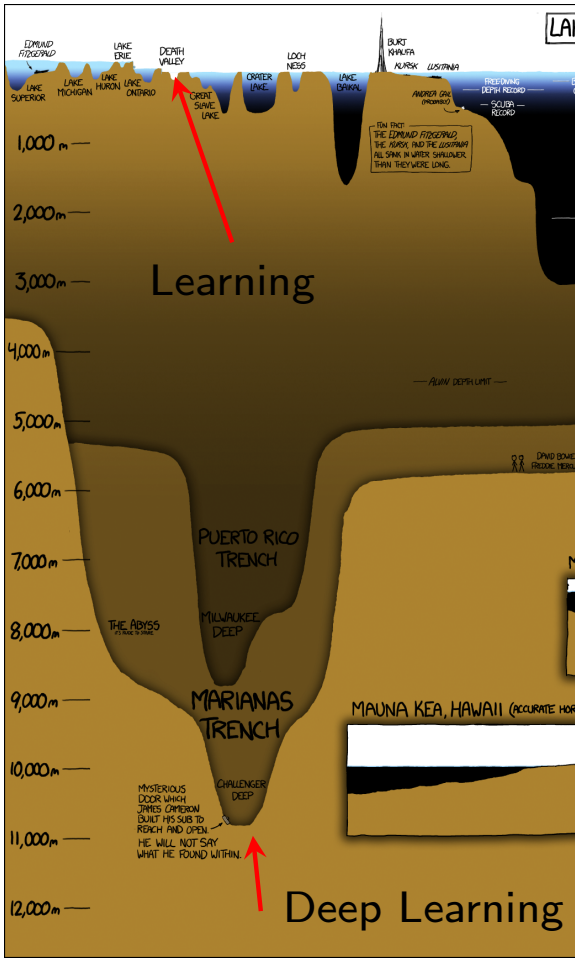

TensorFlow is

- An open source deep learning library

A Tour of TensorFlow

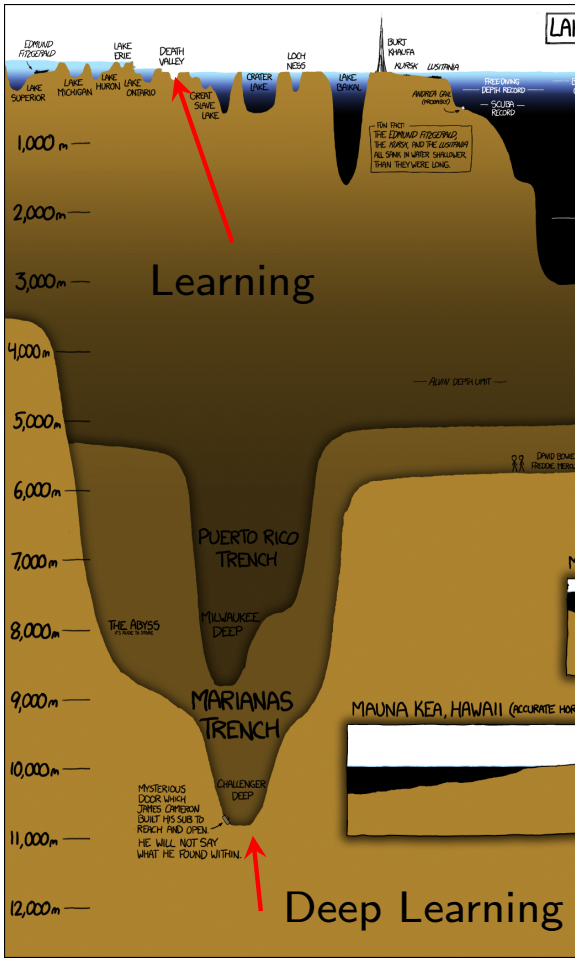

TensorFlow is

- An open source deep learning library
- Released by Google in November 2015

A Tour of TensorFlow

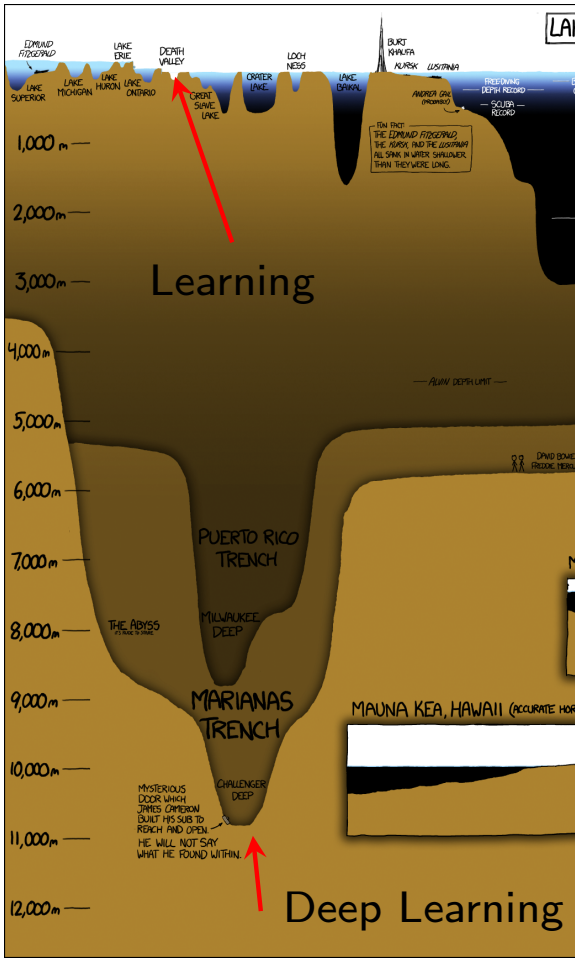

TensorFlow is

- An open source deep learning library
- Released by Google in November 2015
- Especially suited to:

A Tour of TensorFlow

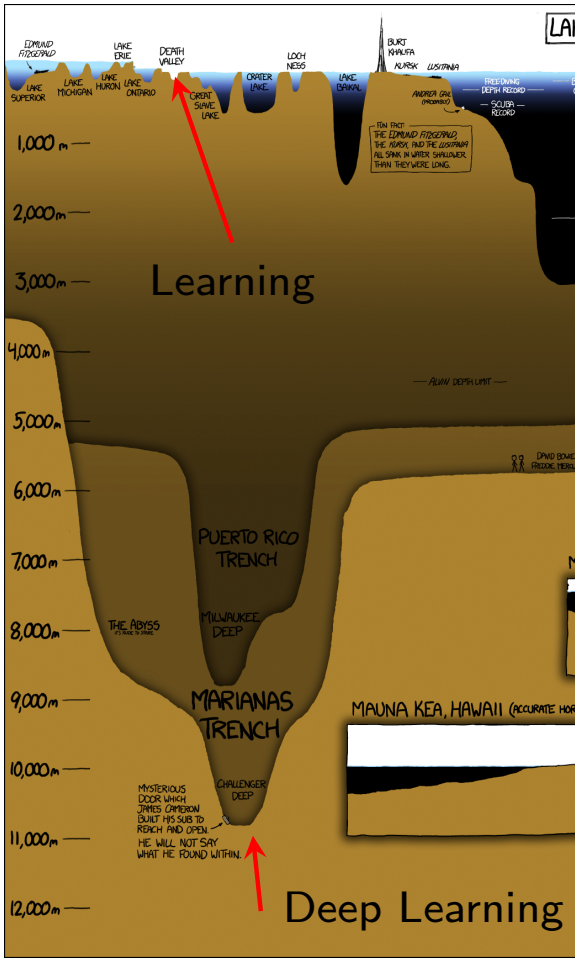

TensorFlow is

- An open source deep learning library
- Released by Google in November 2015
- Especially suited to:
  - "Large-scale machine learning on
  - heterogenous distributed systems"

Contents

Computational Paradigms

Computational Paradigms

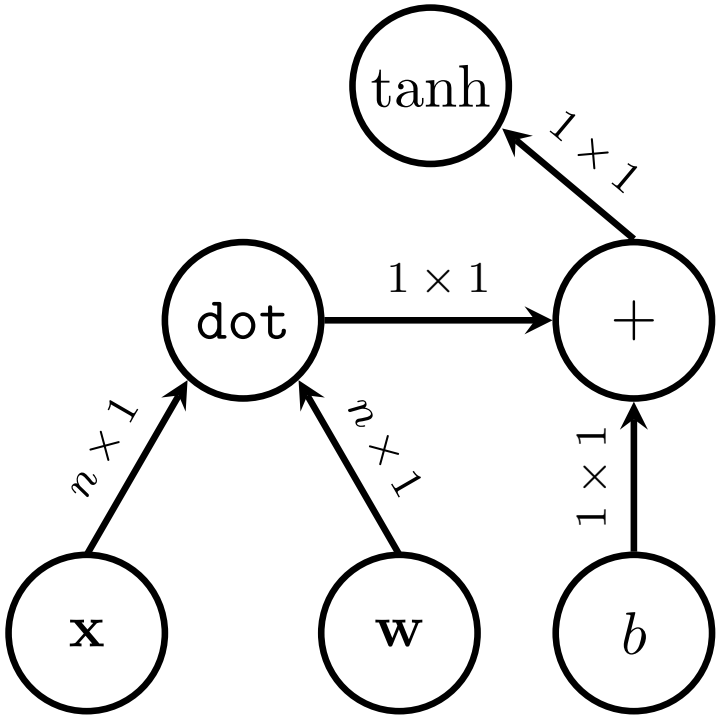

$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

Computational Graphs

Computational Paradigms

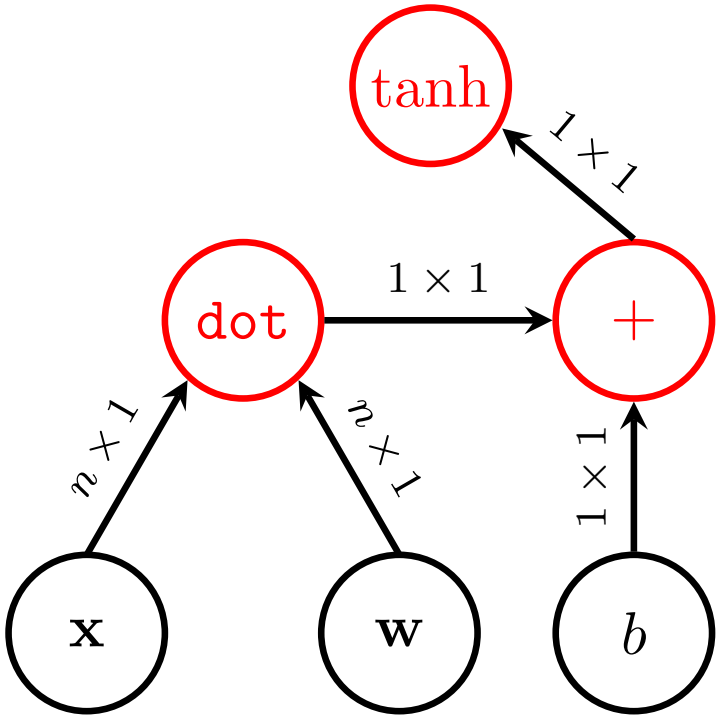

$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

Computational Graphs

1. Operations

Computational Paradigms

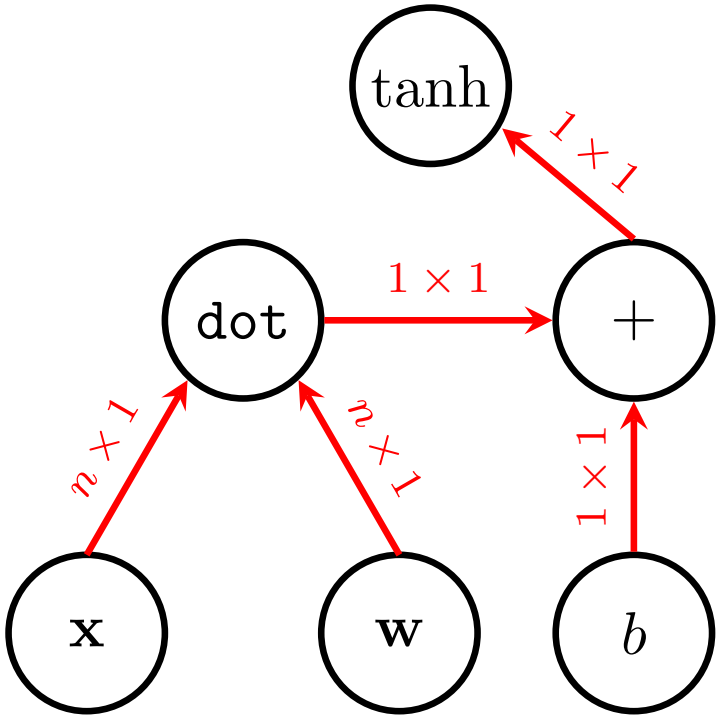

$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

Computational Graphs

1. Operations
2. Tensors

Computational Paradigms

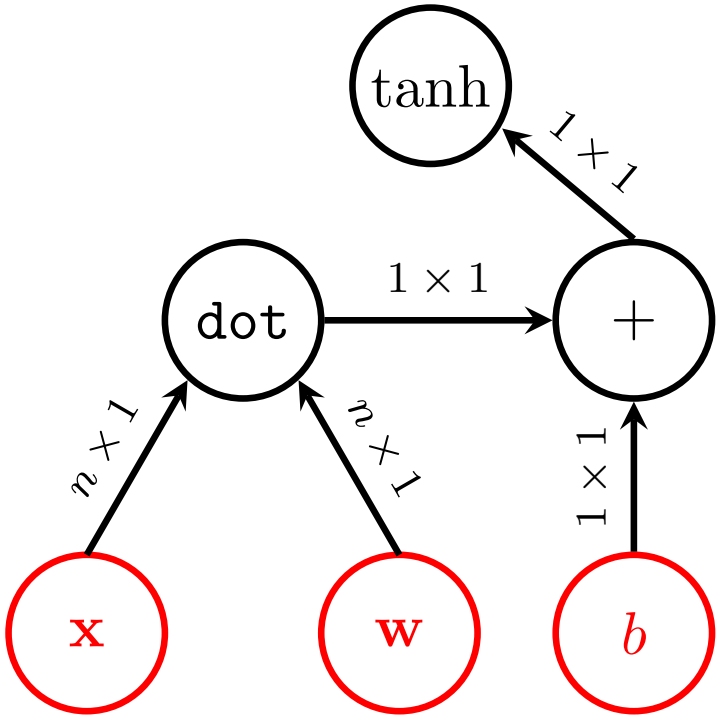

$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

Computational Graphs

1. Operations
2. Tensors
3. Variables

Computational Paradigms

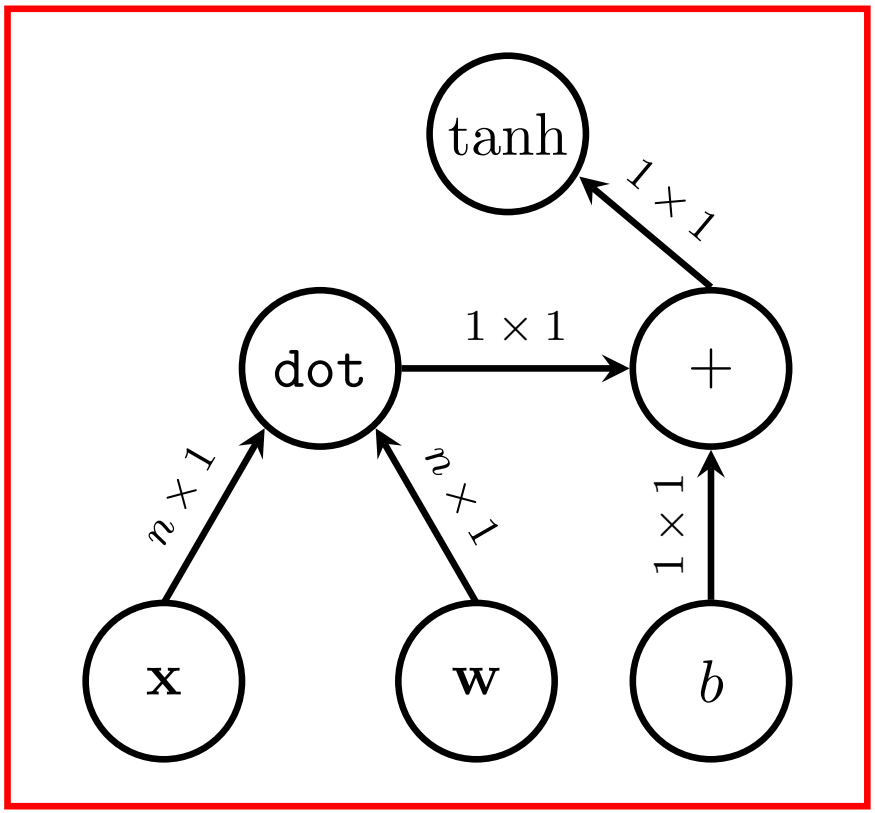

$$\hat{y} = \texttt{session.run}(\tanh(\mathbf{x}^\top \mathbf{w} + b))$$

Computational Graphs

1. Operations
2. Tensors
3. Variables
4. Sessions

Execution Model

Execution Model

Actors

1. Client

Execution Model

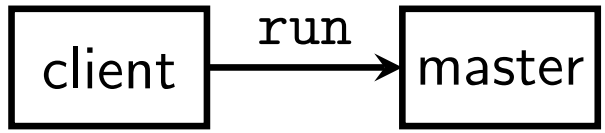

Actors

1. Client
2. Master

Execution Model

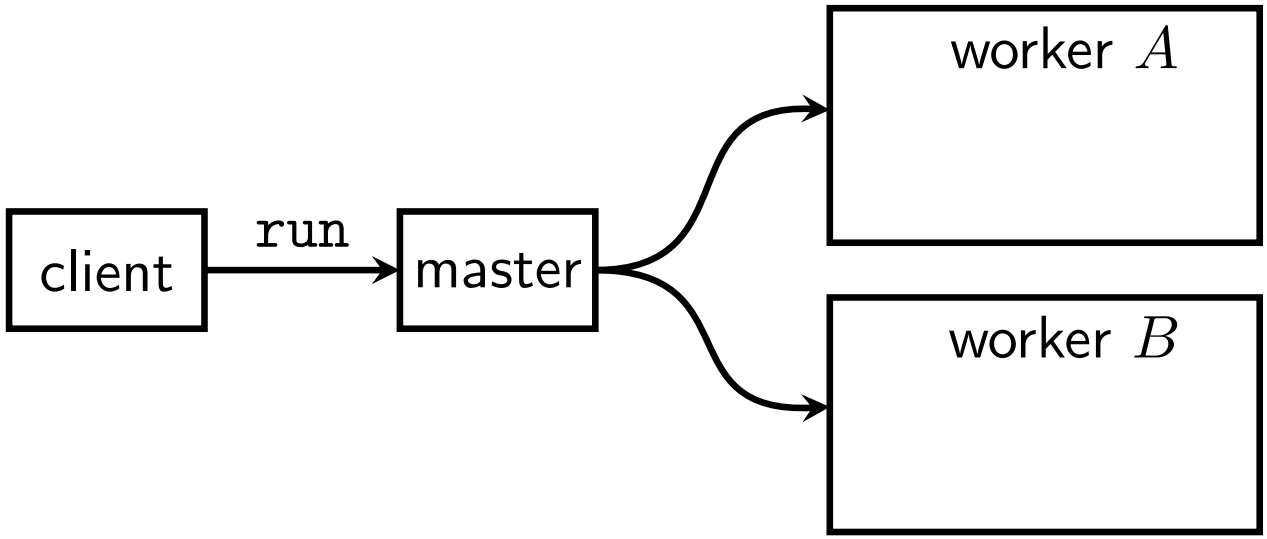

Actors

1. Client
2. Master
3. Workers

Execution Model

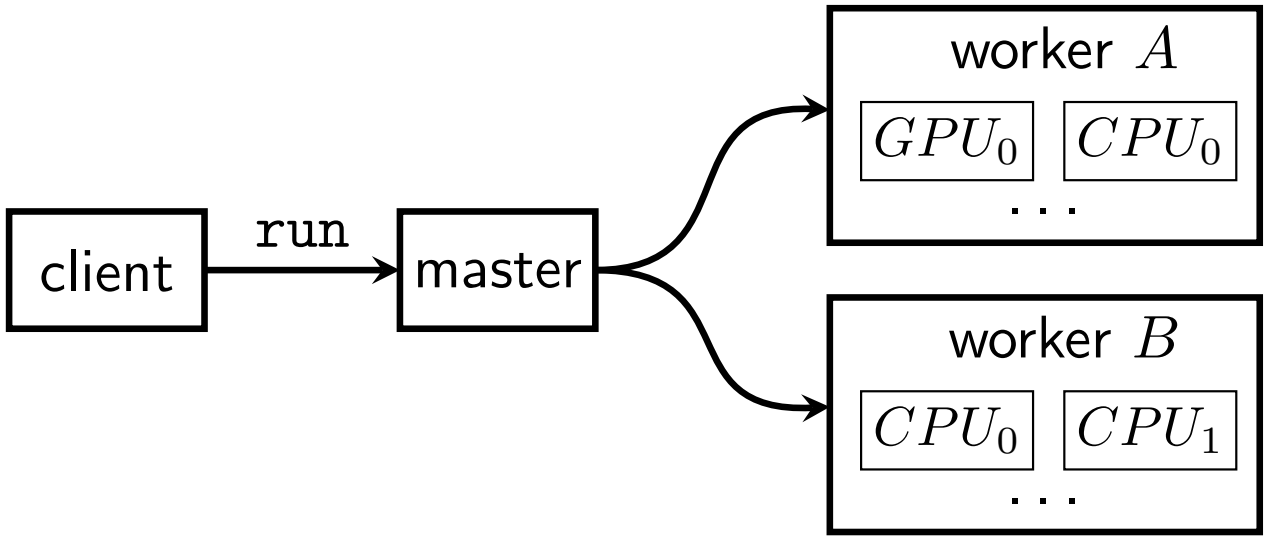

Actors

1. Client
2. Master
3. Workers
4. Devices

Execution Model

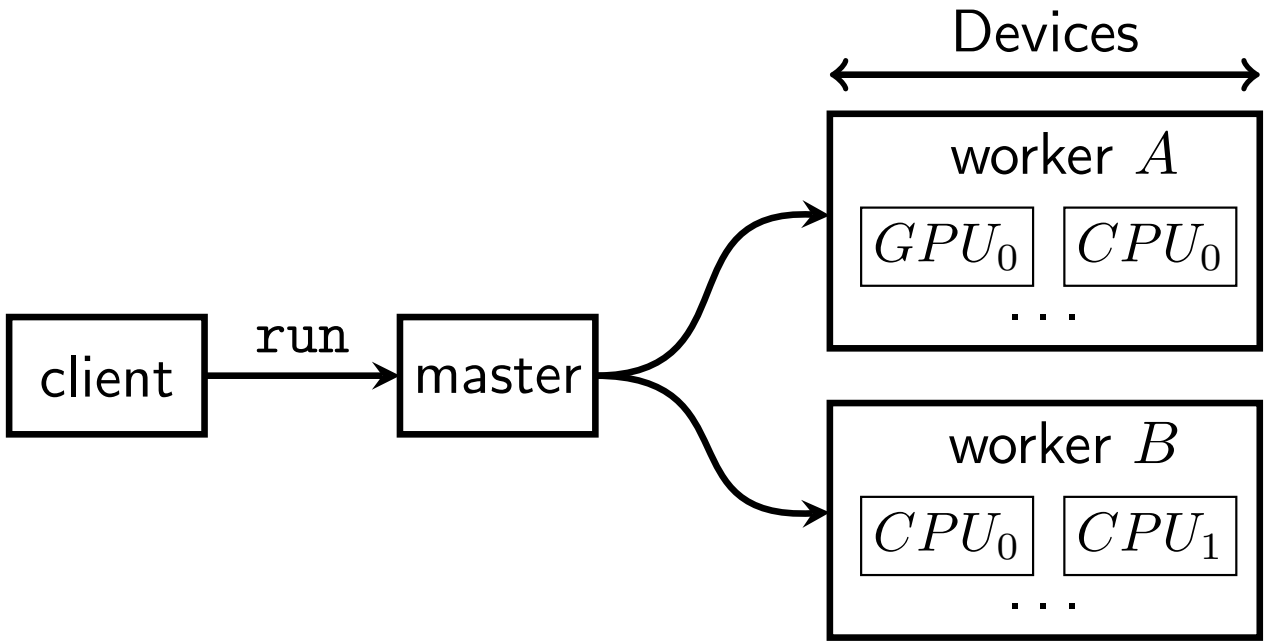

Actors

1. Client
2. Master
3. Workers
4. Devices

Execution Model

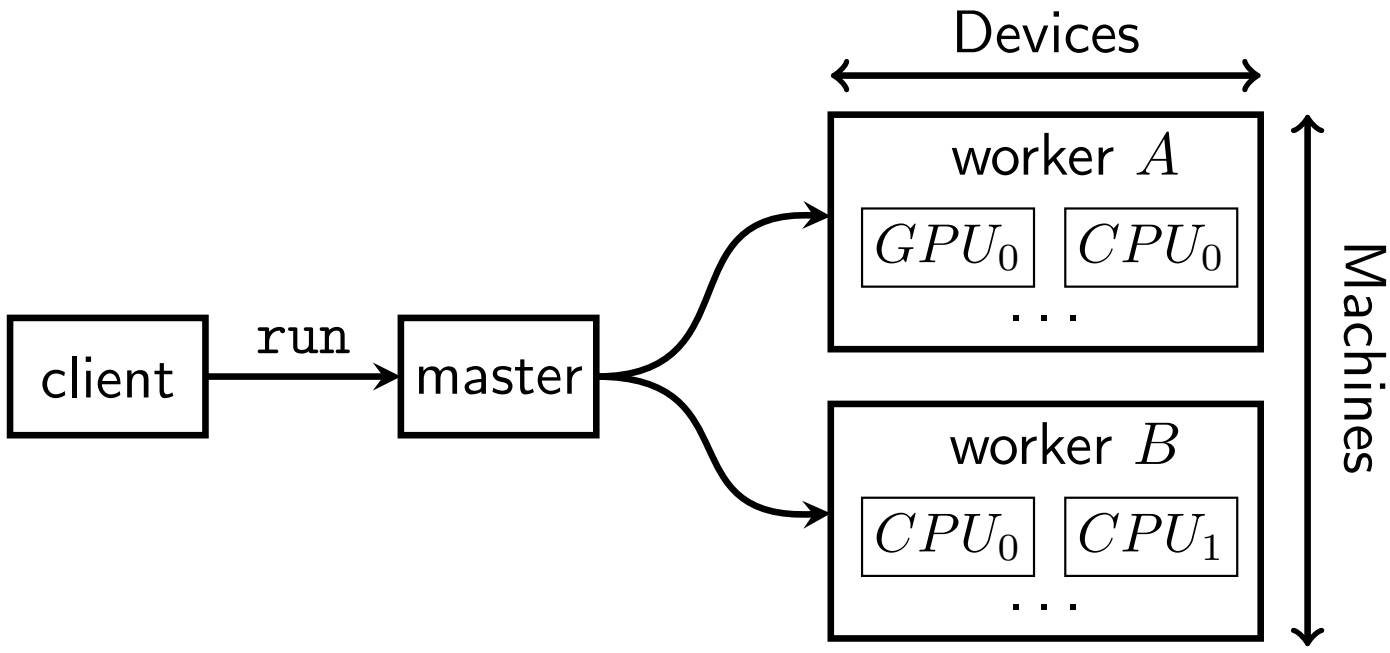

Actors

1. Client
2. Master
3. Workers
4. Devices

Visualization Tools

- Deep Neural Networks have the tendency of being . . . deep
- Easy to drown in the complexity of an architecture
- $>$ 36,000 nodes for Google's *Inception* model

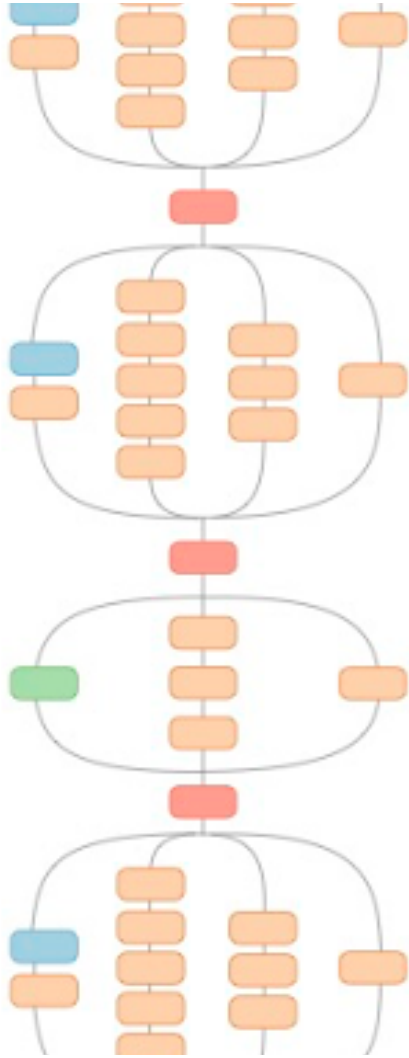

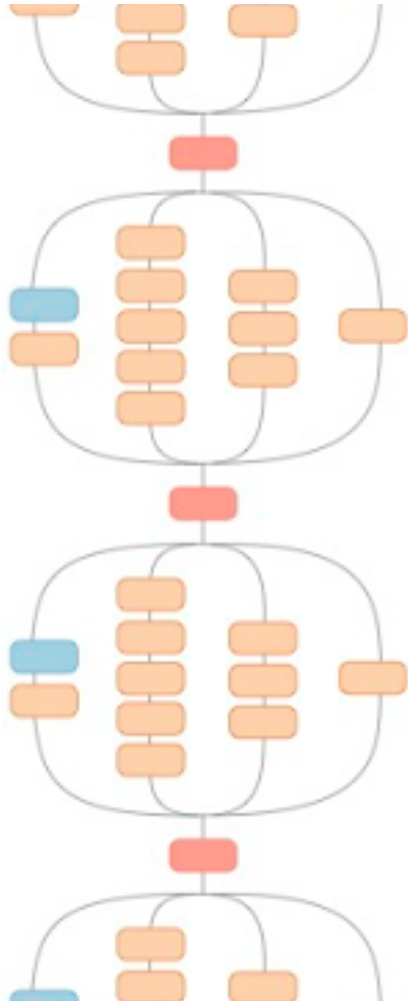

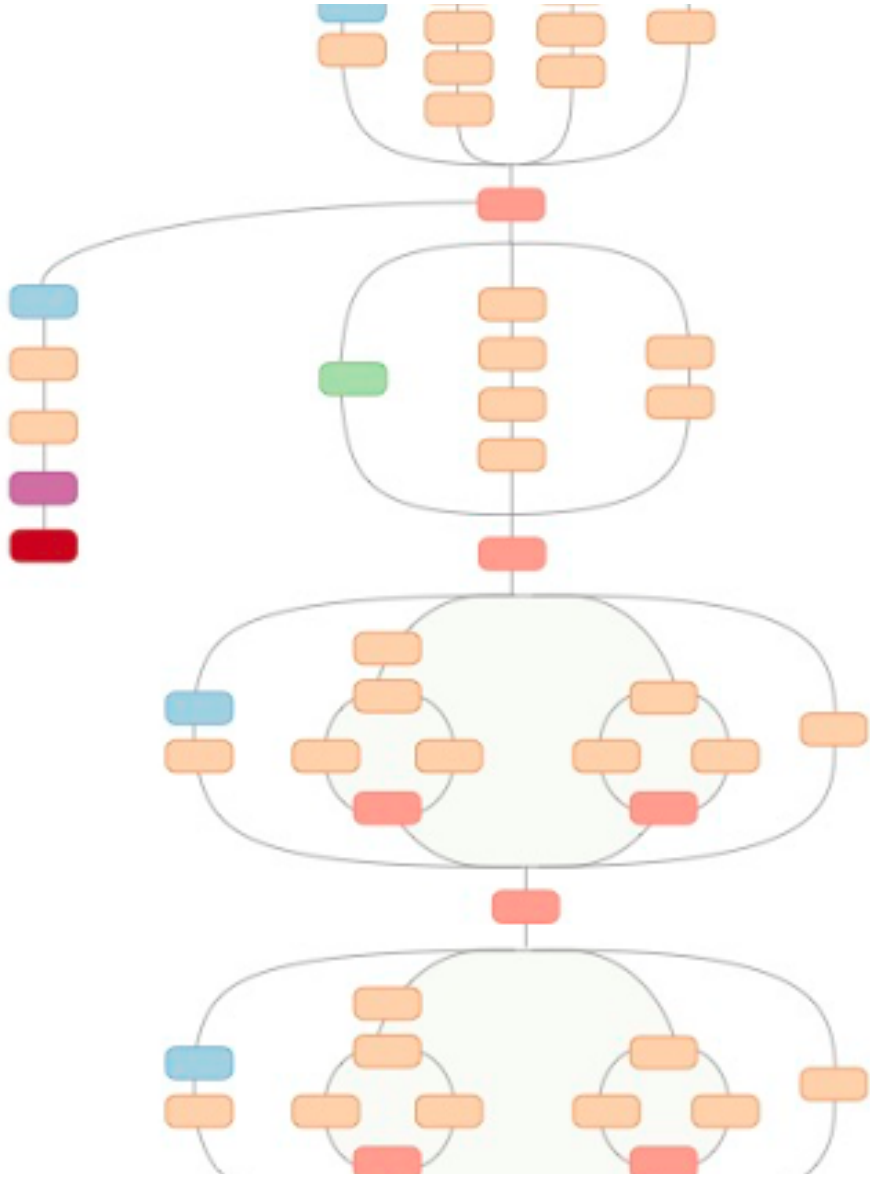

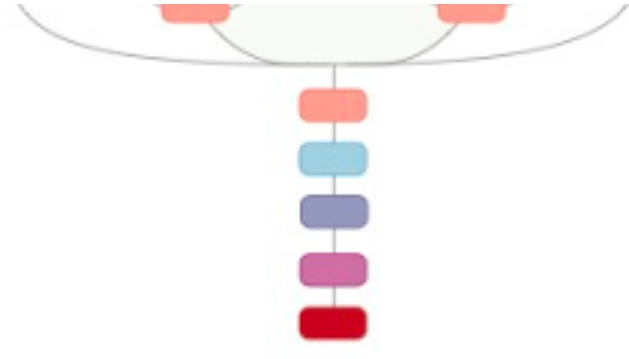

Source: http://googleresearch.blogspot.de/2016/03/train-your-own-image-classifier-with.html

TensorBoard to the Rescue

Use Cases

- Smart email replies in Google *Inbox*
- Emails mapped to "thought vectors"
- LSTMs synthesize valid replies

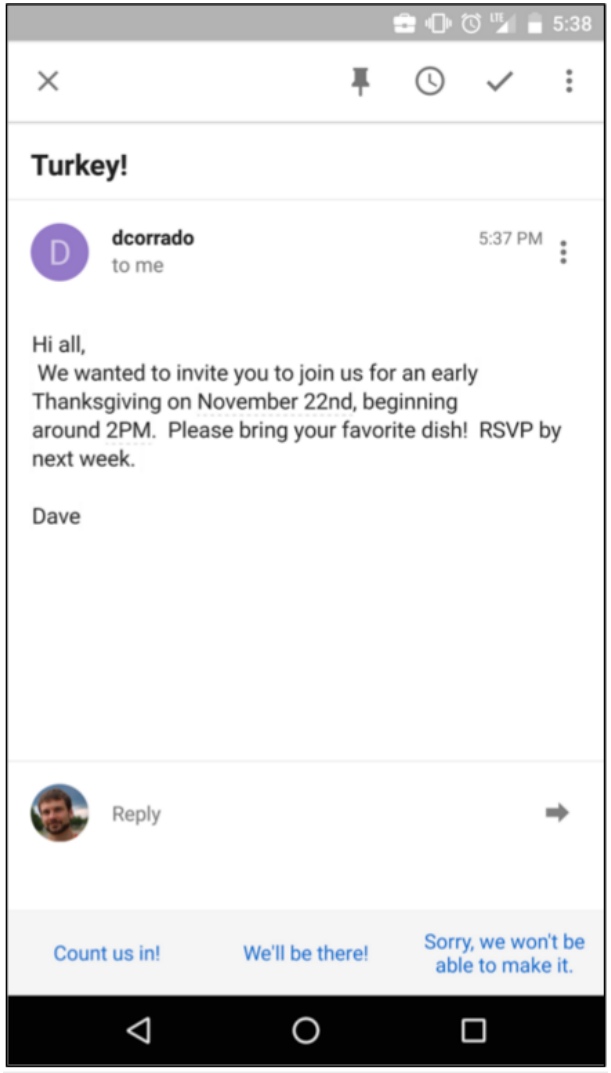

Source: http://googleresearch.blogspot.de/2015/11/computer-respond-to-this-email.html

Use Cases of TensorFlow

- Google DeepMind now using TensorFlow
- Already for *AlphaGo*
- According to a DeepMind SWE reasons are:
  - Integration with Google Cloud Platform,
  - Python,
  - Support for TPUs,
  - Ability to run on many GPUs.

Source: https://deepmind.com/css/images/opengraph/alphago-logo.png

Walkthrough

Thank You